

# LIBXSTREAM Samples

## Ozaki Scheme — OpenCL

High-precision GEMM on OpenCL devices via mantissa slicing (Scheme 1) or Chinese Remainder Theorem (Scheme 2). Both schemes decompose FP matrices into int8/u8 tiles and use DPAS/XXM matrix engines when available. This is an OpenCL adaptation of the CPU-based Ozaki sample in LIBXS.

### Build

```
cd samples/ozaki
make [DBG=1]
```

Requires an OpenCL runtime and headers. BLAS is linked via `BLAS=2` for the reference GEMM.

### Run

```
./ozaki.x [M [N [K [transa [transb [alpha [beta [lda [ldb [ldc]]]]]]]]]]
```

All arguments are positional and optional:

|      |          |         |             |                                                                                                                                                                                                                                                                                                                                                                                                             |
|------|----------|---------|-------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| Pos. | Argument | Default | Description | 1   M   257   Rows of C and op(A)     2   N   M   Columns of C and op(B)     3   K   M   Inner dimension     4   transa   0   0=N, 1=T for A     5   transb   0   0=N, 1=T for B     6   alpha   1   Scalar multiplier for A*B     7   beta   1   Scalar multiplier for C     8   lda   auto   Leading dimension of A     9   ldb   auto   Leading dimension of B     10   ldc   M   Leading dimension of C |
|------|----------|---------|-------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

### Environment Variables

|                         |          |         |             |                                                                                                                                                                                                                                                                                                          |
|-------------------------|----------|---------|-------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Scheme Selection</b> | Variable | Default | Description | OZAKI   3   1=mantissa slicing, 2=CRT, 3=adaptive (default when unset)     OZAKI_FP   64   64=fp64 (double), 32=fp32 (float)     OZAKI_N   (auto)   Slices (Sch.1: fp64=8, fp32=4) or primes (Sch.2: fp64=16, fp32=9)     OZAKI_XOVER   27   Sch.1/2 crossover weight used by OZAKI=3 (fp64, non-NVIDIA) |
|-------------------------|----------|---------|-------------|----------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

OZAKI=3 (adaptive) compares the two costs per call and picks the cheaper path: the Scheme-1 pair count against the Scheme-2 prime count plus its reconstruction, weighted by `OZAKI_XOVER` and amortized over K. The comparison is stateless — it uses the static slice cutoff, not a measured one — so mixed matrix sizes in one process cannot influence each other. On NVIDIA it selects Scheme 2 outright, and in fp32 on other devices Scheme 1, both because counting GEMMs mispredicts there.

|                 |          |         |             |                                                                                                                                                                                                                                                                                                                                                                                                                                           |
|-----------------|----------|---------|-------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|
| <b>Accuracy</b> | Variable | Default | Description | OZAKI_FLAGS   3   Sch.1 bitmask: 1=Triangular, 2=Symmetrize, 0=full S^2. No Sch.2     OZAKI_TRIM   0   Precision levels to trim (0=exact). ~7 bits (Sch.1), ~4 bits (Sch.2)     OZAKI_I8   0   Sch.2: use signed i8 residues (moduli<=128) instead of u8     OZAKI_GROUPS   0   Sch.2: K-grouping factor, consecutive K panels share reconstr.     OZAKI_FRACCRT   (auto)   Sch.2: 0=Garner, 2=fractional CRT. Auto: 0 if unfused, else 2 |
|-----------------|----------|---------|-------------|-------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------------|

`OZAKI_FLAGS` selects how the slice-pair loop is traversed, and the two bits mean the same on the host and on the device: 1 starts the inner slice at the outer one, 2 adds the transposed product of every off-diagonal pair. Only 0 (the full square loop) and 3 (both bits) compute the whole sum – 1 alone omits the transposed products and 2 alone counts them twice, so both are for symmetric operands only.

`OZAKI_GROUPS` splits K into panels to bound the preprocessing footprint, and it is a slowdown at every size measured because it is incompatible with the unfused epilogue (`OZAKI_UNFUSE`, on by default for GPUs) – use it only when the per-pass buffers do not fit.

`OZAKI_FRACCRT=1` trades exactness for speed (flat fractional sum, magnitude-bounded); modes 0 and 2 are both exact and differ only in speed. Which one is faster depends on the epilogue rather than on the device: with `OZAKI_UNFUSE` the reconstruction is a separate pass and Garner wins, without it the fractional variant does, so the default follows that knob.

**Hardware Control** | Variable | Default | Description | |||| | OZAKI\_TM | (auto) | Output tile M (BM). Overrides size-aware selection | | OZAKI\_TN | (auto) | Output tile N (BN). Overrides size-aware selection | | OZAKI\_RTM | (auto) | Register tiling M (power of two). Auto: 2 (HIER), 4 (256-GRF) | | OZAKI\_RTN | (auto) | Register tiling N. Auto: 8 (NV MMA Sch.2), 2 (Intel), 1 (other) | | OZAKI\_SB | 1 | Sch.1: slice-block width for the pair loop (needs 256-GRF) | | OZAKI\_WGS\_MAX | (auto) | Work-group size ceiling for tile selection (0=hardware bound) | | OZAKI\_LU | 0 | Unroll hint for the K-loop (0=compiler default) | | OZAKI\_SKIP\_GARNER | 0 | Sch.2: skip reconstruction (timing only, result is not a GEMM) | | OZAKI\_WG | 0 | Work-group size hint (0=no hint) | | OZAKI\_SG | (auto) | Sub-group size (forced to 16 with XMX) | | OZAKI\_BIGGRF | (auto) | Override 256-GRF detection (0=off, 1=on). HIER defaults to 128 | | OZAKI\_KU | 2 | K-loop unroll factor (16 with OZAKI\_WGMMA, 8 if RS is off) | | OZAKI\_RC | 8 | DPAS repeat count (8 or 4) | | OZAKI\_PB | 1 | Sch.2: CRT prime batching factor | | OZAKI\_HIER | (auto) | Sch.2: hierarchical CRT (default on). Two-level Garner reconstr. | | OZAKI\_PREFETCH | 0 | Sch.1: enable prefetching | | OZAKI\_SCALAR\_ACC | 0 | Sch.1: force scalar accumulation | | OZAKI\_BVNNI | 1 | NVIDIA: pre-interleave B so each operand is one aligned uint | | OZAKI\_BKMAJOR | 0 | NVIDIA, Sch.2: transpose B to [N][K] instead (see below) | | OZAKI\_BBLOCK | (auto) | Sch.2: block 16 K-values of a B column. On with wgmma (below) | | OZAKI\_WGMMA | (auto) | Sch.2: warp-group MMA. On where reachable (see below) | | OZAKI\_WGMMA\_N | 128 | Warp-group tile width, 64 or 128 | | OZAKI\_WGMMA\_M | 128 | Warp-group tile rows: 128 = two warp groups, 64 = one | | OZAKI\_WGMMA\_RS | 1 | Warp-group MMA takes A from registers instead of shared memory | | OZAKI\_UNFUSE | (auto) | Sch.2: reconstruct in a 2nd kernel. On for GPUs (see below) | | OZAKI\_SWIZZLE | 0 | Sch.2: work-group rasterization width (0=launch order) | | OZAKI\_arena | (auto) | Device scratch arena in MB (0=off). Auto: on if no pool |

On NVIDIA GPUs the Scheme-2 default RTN=8 is tuned for large K: it doubles the output tile per work-group, which needs a long K-loop to pay for itself (+38% at n=4096). Below K of about 1024 it costs 6% (K=512) to 21% (K=256) instead — set OZAKI\_RTN=4 there.

Warp-group MMA runs the Scheme-2 GEMM on Hopper's `wgmma` instead of `mma.sync`. It is the default wherever it is reachable, and OZAKI\_WGMMA=0 selects `mma.sync`. On an H100 PCIe it is faster at every size and bit-identical: the GEMM and reconstruction rows together are 13.66 -> 6.12 ms at n=4096 and 101.8 -> 45.8 at n=8192 against `mma.sync` at its own defaults, and `transa`, `transb`, rectangular and non-tile-multiple shapes gain the same. It needs hierarchical CRT, no K-grouping and OZAKI\_PB=1, and it fixes the tile to OZAKI\_WGMMA\_MxOZAKI\_WGMMA\_N regardless of OZAKI\_TM/OZAKI\_TN.

“Reachable” is established by building a throwaway kernel during initialization rather than by inspecting the device, because nothing the device reports settles it: the advertised local-memory size is 48 KB on a part that accepts the 128 KB this path uses, and compute capability 9.0 covers hardware where the instruction no longer exists. If that build fails, initialization says so under OZAKI\_VERBOSE=1 and the `mma.sync` path is used with its own tuning, at no cost in speed.

Three knobs tune it, all set to the fastest measured value by default. OZAKI\_WGMMA\_RS=0 stages A in shared memory instead of loading it into registers, and costs 10% (n=1024) to 40% (n=8192). OZAKI\_KU sets how much K is staged per round (16 here, or 8 without RS); it is what the path spends shared memory on, 128 KB at the defaults, so lower it if a device refuses that. OZAKI\_WGMMA\_M=128 runs two warp groups per work-group instead of one, worth +4% at n=2048 and +18% at n=8192.

Two implications worth knowing. The kernel is built twice — once from OpenCL C, then again from patched PTX — because warp-group MMA cannot be expressed in OpenCL C, so the first JIT of a shape costs about double. And if that patching fails the kernel is refused rather than run, since an unpatched kernel would silently compute zeros; OZAKI\_VERBOSE=1 reports it.

OZAKI\_UNFUSE=1 moves the CRT reconstruction out of the GEMM into a second kernel: the GEMM stores one residue byte per prime and output, `gemm_crt_reduce` reads them back and reconstructs C. It is the default on GPUs, OZAKI\_UNFUSE=0 selects the fused epilogue, and both produce identical results.

This is faster because of what it removes rather than what it adds. With the prime loop outermost, the fused kernel must keep every output's group values live across it, which is 2 KB per work-item of dynamically indexed arrays — more than the L1 can hold for a work-group, so the epilogue runs at memory speed. A separate pass puts the output loop outermost and needs only four group values, i.e. registers. Measured at n=4096: 13.08 ms fused against 7.87 + 0.66 unfused (+53%), and the GEMM alone beats even its own loop-only time, because the frame was slowing the K-loop too. The gain is largest where the epilogue dominated: +158% at n=257, +80% at n=2048, +76% in fp32. It is not warp-group specific and helps `mma.sync` by 11%. Nor is it NVIDIA-specific: on a PVC (DPAS) the same figures are 0.66 -> 0.28 ms at n=257, 26.5 -> 22.0 at n=4096 and 15.8 -> 12.6 in fp32, also bit-identical.

The residue planes it needs are part of the per-call device scratch, together with the operand planes and the device copy of C. OZAKI\_arena carves all of it from one allocation kept for the life of the context instead of creating and destroying

it per call, which is worth a factor of 6.7 at  $n=4096$  on a GH200 and 18% on an H100 PCIe. It is enabled by default wherever libxstream has no memory pool, so on NVIDIA nothing needs setting.

`OZAKI_ARENA=0` restores per-call allocation. A positive value caps the arena in MB, and sizing it wrong is worse than not capping: a call needing more than the cap falls back per buffer, and a first call that exceeds it gets no arena at all. The requirement is roughly 1.2 GB at  $n=4096$  in FP64 and scales with  $n^2$ , or call `ozaki_scratch_size`.

Set a positive value explicitly on a device that does have a pool but still wants the arena - `LIBXSTREAM_USM` on NVIDIA is that case, since the pool alone turns the arena off. Planes kept by `OZAKI_CACHE` are never carved, so the two knobs are independent. A caller that owns device memory can supply it instead with `ozaki_scratch_size/ozaki_scratch_set` (the benchmark does so under `OZAKI_SCRATCH=1`); the interceptor in LIBXS relies on the default, since a BLAS call cannot be handed a workspace.

Two things to know. Reading the profile, the two kernel rows have to be added for a total; the FLOP rate is attributed to the GEMM row alone. And it needs scratch memory of `OZAKI_N` bytes per output element, twice the size of `C` in fp64, which is freed per call.

`OZAKI_BBLOCK` stores `B` so that 16 consecutive `K`-values of one column are contiguous, with columns 16 bytes apart. It is the default wherever warp-group MMA runs, and `OZAKI_BBLOCK=0` selects the 4-byte interleave.

The reason is copy count. The warp-group loop is bound by how many `cp.async` instructions the staging needs, not by bandwidth: splitting `A`'s 16-byte copies into 4-byte ones costs 107% for the same bytes. `B` staged from the interleave needs four times the copies of `A`, and this layout removes three quarters of them while keeping both sides coalesced - the transposed layout gives the consumer its 16 bytes but scatters the producer, and the interleave coalesces both but forces the 4-byte copies. Measured: the GEMM gains 18% at  $n=4096$  (8.03 -> 6.89 ms), 14% at  $n=8192$ , 55% at  $n=1024$  and 40% at  $n=257$ , while `B` preprocessing pays 0.53 -> 0.73 ms, so about 11% net at  $n=4096$  and more below it. Bit-identical, and the lane must walk columns rather than `K`-blocks - mapping it the other way reads 64 KB apart and loses 17% instead.

`OZAKI_SWIZZLE=W` walks the tile grid in strips `W` tiles wide instead of in the launch order, so the work-groups resident at one time cover a block rather than a column of the grid and re-read the residue planes fewer times. It is off by default because it barely pays: measured at  $n=8192$  it gains 4% on the GEMM (68.9 -> 66.1 ms at `W=8`) and costs the reconstruction pass 13% (2.51 -> 2.57), for about 4% net, and at  $n=4096$  it is a wash. Worth trying only for large square problems.

`OZAKI_BKMAJOR=1` selects the transposed `B` layout, which exists for warp-group MMA (both operands `K`-major) and is exact like the default but slower for the kernels shipping today: at  $n=4096$  it costs the fused GEMM 15.4 -> 21.0 ms and `B` preprocessing 0.5 -> 3.5 ms, because consecutive lanes then read `k_pad` apart instead of 4 bytes apart. It does not pay for the warp-group path either, although that one stages `B` cooperatively and could use 16-byte copies from a `K`-major layout: it measured 8.1 -> 9.5 ms at  $n=4096$ , because the loss in global-side coalescing outweighs the cheaper staging.

The hierarchical CRT groups primes for the two-level reconstruction, and the group size is derived from the prime count rather than fixed: 4 is the maximum the 32-bit level-2 datapath allows, but a group left holding a single prime is pathological. At the fp32 count of 9 primes, groups of 4 leave a 4,4,1 split whose reconstruction costs 1.21 ms at  $n=4096$  - twice what a 12-prime run of three full groups costs - so fp32 uses groups of 3 and measures 0.83 ms. fp64 (16 primes) keeps groups of 4 and is unaffected. Regrouping changes the order of the final Horner summation, so an fp32 result can differ in the last bits from an earlier release, at the same accuracy (`1inf_re1` around  $1e-06$  either way); integer CRT reconstruction is exact regardless of grouping.

No other setting differs by precision: the tile, staging depth, `B` layout and warp-group geometry were swept for fp32 and rank exactly as in fp64.

**Memory and Caching** | Variable | Default | Description | |||| | `OZAKI_CACHE` | 0 | Preprocessing cache bitmask: 1=A, 2=B, 3=both (or any negative) | | `OZAKI_NPANEL` | 1 | Sch.2: N-panel width (0=auto, 1=disable). See Panel Pipeline |

The preprocessing cache also stores the last effective cutoff from Scheme 1 occupancy detection. On cache hits the cutoff is reused without device-to-host readback, eliminating the sync bubble.

|                  |  |          |  |                                                  |  |             |  |  |  |         |  |   |  |                                 |  |  |
|------------------|--|----------|--|--------------------------------------------------|--|-------------|--|--|--|---------|--|---|--|---------------------------------|--|--|
| <b>Benchmark</b> |  | Variable |  | Default                                          |  | Description |  |  |  | NREPEAT |  | 1 |  | Number of benchmark repetitions |  |  |
| OZAKI_VERBOSE    |  | 0        |  | 0=silent, 1=errors, 2=warnings, 3+=all. Neg.=all |  |             |  |  |  |         |  |   |  |                                 |  |  |

Additional variables for accuracy monitoring and complex GEMM dispatch (OZAKI\_THRESHOLD, OZAKI\_STAT, OZAKI\_EPS, OZAKI\_RSQ, OZAKI\_EXIT, OZAKI\_COMPLEX) are handled by the LIBXS Ozaki sample (LIBXS), which owns the GEMM interceptor. See its README for details.

Kernel timings come from LIBXSTREAM rather than from a sample-specific facility: LIBXSTREAM\_PROFILE=1 reports one row per kernel (preprocess A, preprocess B, the fused GEMM, and so on), and LIBXSTREAM\_PROFILE\_MEM=1 reports transfer rates. Neither needs a phase to be selected up front — every kernel is recorded and the interesting rows are read from the report.

**cuBLAS Reference** When the OpenCL headers are taken from a CUDA installation, the sample additionally links cuBLAS and reports a device-side reference GEMM (cuBLAS GEMM and cuBLAS DIFF). Build with make CUBLAS=0 to opt out. The host BLAS stays the accuracy reference for both the Ozaki and the cuBLAS result, and a failing cuBLAS run is not fatal.

By default the sample requests FP64 emulation without the built-in fallback to native FP64 (CUDA 13.0u2 and later, compute capability 8.0 and later). The variables below are populated before cuBLAS is entered, hence any of them that is already set wins:

|  |          |  |                |  |                                              |  |                                 |  |   |  |  |                           |  |   |  |  |
|--|----------|--|----------------|--|----------------------------------------------|--|---------------------------------|--|---|--|--|---------------------------|--|---|--|--|
|  | Variable |  | Sample default |  |                                              |  | CUBLAS_EMULATE_DOUBLE_PRECISION |  | 1 |  |  | CUBLAS_EMULATION_STRATEGY |  | 0 |  |  |
|  | eager    |  |                |  | CUBLAS_EMULATION_SPECIAL_VALUES_SUPPORT_MASK |  | 0                               |  |   |  |  |                           |  |   |  |  |

Setting CUBLAS\_EMULATE\_DOUBLE\_PRECISION=0 yields the native device GEMM, which is the baseline to compare against. For OZAKI\_FP=32 the counterpart is CUBLAS\_EMULATE\_SINGLE\_PRECISION, which requires compute capability 10.0 and later.

|  |          |  |         |  |             |  |  |  |                   |  |   |  |                                                       |  |  |                   |  |   |  |                                                          |  |  |
|--|----------|--|---------|--|-------------|--|--|--|-------------------|--|---|--|-------------------------------------------------------|--|--|-------------------|--|---|--|----------------------------------------------------------|--|--|
|  | Variable |  | Default |  | Description |  |  |  | OZAKI_CUBLAS_BITS |  | 0 |  | Mantissa bits: 0=default, <0=match the OZAKI_N slices |  |  | OZAKI_CUBLAS_XPTR |  | 0 |  | 1=pass LIBXSTREAM device pointers to cuBLAS (experiment) |  |  |
|--|----------|--|---------|--|-------------|--|--|--|-------------------|--|---|--|-------------------------------------------------------|--|--|-------------------|--|---|--|----------------------------------------------------------|--|--|

A non-zero OZAKI\_CUBLAS\_BITS fixes the number of int8 slices (`slice = ceil((bits + 1) / 8)`) instead of letting cuBLAS pick the precision per call. A negative value matches the component count of this sample, which is a slice count for OZAKI=1 but a prime count for OZAKI=2 — only the former is comparable.

OZAKI\_CUBLAS\_XPTR=1 is an experiment and expected to fail: the device-side pointers of LIBXSTREAM are not addresses of the CUDA context that cuBLAS runs in.

Like the Ozaki timing, the cuBLAS timing covers the transfers, i.e. A, B and C are uploaded and C is read back per iteration; CUDA events report the device-side split (`gemm`, `h2d`, `d2h`) separately. With OZAKI\_CUBLAS\_XPTR=1 the transfers are not on the CUDA timeline and read as zero.

The host buffers are pinned by OpenCL rather than by CUDA. LIBXSTREAM registers its host allocations with the CUDA runtime whenever the application links it, so no action is needed here; LIBXSTREAM\_VERBOSE=2 reports how many were accepted. LIBXSTREAM\_PIN=0 leaves them unregistered, which makes the cuBLAS transfers pageable — several times slower, and no longer comparable to the Ozaki row. LIBXSTREAM\_PIN=3 loads the CUDA runtime when the application did not link it, which is what a wrapper driver needs: it links libOpenCL alone, so the default setting finds no entry point and registers nothing.

Two properties to keep in mind when reading the numbers. The mode printed alongside the timing is the one that was requested: whether a call was emulated cannot be queried, and `nsys profile ./ozaki.x` is the way to confirm it. And emulation needs a workspace, which the sample supplies (`-DOZAKI_CUBLAS_WORKSPACE=<bytes>`, 2 GB by default, 0 to disable) because a workspace that is too small does not fail the call but silently falls back to non-emulated kernels.

### Output Tile Selection

The output tile per work-group (BM x BN) is chosen per call from M and N, because the best tile depends strongly on problem size: the total work-item count is invariant to the tile, so the tile only controls how many work-groups the problem splits into and how far M and N are padded to a tile multiple. A large tile maximizes operand reuse but produces too few work-groups to fill the device at small sizes; on a GPU Max 1550 the fixed 128x256 tile costs up to 1.7x at M=N=256..640 while being optimal from ~1024 upward.

Selection maximizes arithmetic intensity  $BM \cdot BN / (BM + BN)$  per unit of padded work, subject to two constraints: the work-group size bound  $SG * (BM / (XMX\_M * RTM)) * (BN / (XMX\_N * RTN)) \leq \text{max\_wgs}$ , and enough work-groups to saturate the compute units (OZAKI\_TILE\_SAT). The intensity term is maximal for square tiles, which keeps the split balanced rather than degenerate. Set OZAKI\_TM/OZAKI\_TN to bypass selection and pin a tile.

The register tiling is selected per call as well, on Intel GPUs: the auto value in the table is the small-problem choice, and a larger one is taken once the problem still forms one tile per compute unit at the coarser granularity that implies. OZAKI\_VERBOSE=3 prints it as `rtm=2..4`, meaning either may be used, and the per-call value appears in the JIT line as `rt=`. Setting OZAKI\_RTM or OZAKI\_RTN pins both and bypasses this.

## Panel Pipeline (Scheme 2)

Only ~30% of a Scheme-2 call is the GEMM kernel; the rest is uploading A/B, preprocessing them into CRT residues, and downloading C. Scheme 2 therefore splits **N** into panels and pipelines them: while panel *j* runs its GEMM, panel *j+1* uploads and preprocesses its B columns and panel *j-1* downloads its C block. A is uploaded and preprocessed once as a prologue.

**N** is the right axis because each panel then owns a disjoint column block of B and C. C is read and written exactly once no matter how many panels there are, and panel GEMMs are mutually independent. Splitting K instead would re-read and re-write all of C per group — measured at several ms per pass at `n=4096`, which is most of what the pipeline is trying to hide — and would serialize the panel GEMMs through that accumulation.

Panel B-slices are double-buffered across OZAKI\_NSLOTS slots (2 by default; 3 and 4 measured slower). A panel reusing a slot waits only on the GEMM that last read it, which is what decouples consecutive panels. Because a panel is a slice of B rather than all of B, panelling also lowers peak memory relative to the unpanelled path.

Panelling is skipped when it cannot pay: too few tiles per panel to saturate the device,  $N \leq tn$ , K-grouping active, device-resident operands, or OZAKI\_NPANEL=1. Requesting a B cache (OZAKI\_CACHE=2 or 3) takes precedence — a cached slice buffer must hold all of B, so caching and panelling are alternative ways to avoid the same work (reuse across calls versus overlap within one). With `transb` a panel is a strided row block, so B is uploaded whole and only the preprocessing and download overlap.

Disabled by default because the automatic width does not account for shape: measured on PVC (fp64) it costs up to 12% on square shapes and gains up to 7% where **N** is much larger than **M**. Set OZAKI\_NPANEL=0 to enable it where **N** is the dominant dimension.

## Kernel Registry

Both schemes compile fused GEMM kernels on demand via a JIT registry keyed by the output tile plus the bounds-checking variant; Scheme 1 additionally keys on the compile-time cutoff (OZAKI\_CUTOFF), which lets the compiler eliminate dead slice-pair iterations and reduce register pressure. The first call with a given key triggers JIT compilation (~100 ms); subsequent calls hit the registry cache. A workload with a single matrix shape produces one specialization per scheme.

## Example

```
./ozaki.x 256
```

Scheme 2 on a large matrix:

```
OZAKI=2 ./ozaki.x 4096
```

Adaptive scheme selection with caching:

```
OZAKI=3 OZAKI_CACHE=3 ./ozaki.x 4096
```

## Quick Tuning Guide

Scheme 2 (CRT, OZAKI=2, default): fixed cost of P integer GEMMs plus hierarchical Garner reconstruction. Predictable performance regardless of data distribution. Use OZAKI\_GROUPS for K-grouping at large sizes. The hierarchical CRT (OZAKI\_HIER, on by default) halves private residue arrays and enables GRF128 for doubled thread occupancy. Selecting flat reconstruction takes OZAKI\_HIER=0 OZAKI\_FRACCRT=0 with OZAKI=1 or OZAKI=2, because the per-group fractional CRT and the adaptive scheme both need the hierarchy; the combination warns when it cannot be honoured.

Scheme 1 (mantissa slicing, OZAKI=1): up to  $S*(S+1)/2$  integer GEMMs, but adaptive cutoff can reduce this substantially for narrow exponent spans. Use OZAKI\_TRIM to trade accuracy for speed.

Adaptive (OZAKI=3): automatically picks the cheaper scheme per call based on preprocessing occupancy. Best with OZAKI\_CACHE=3 to avoid repeated occupancy readbacks.

Enable OZAKI\_CACHE=3 when A or B stays constant across calls. Note that only OZAKI\_CACHE=0 disables the cache: a negative value is another spelling of 3.

## Small Matrix Multiplication (SMM) — OpenCL

Batched small matrix multiplications (SMM) on OpenCL devices via the ACC LIBSMM interface. Originated from the DBCSR OpenCL backend, adapted to run on top of LIBXSTREAM. Includes GPU-accelerated kernels, a benchmark driver, and an auto-tuning framework.

### Build

```
cd samples/smm
make [WITH_GPU=<device>] [ELEM_TYPE=<type>]
```

Produces acc\_bench.x. Requires an OpenCL runtime, BLAS (BLAS=2), and LIBXS built from a sibling directory.

| Make variable | Default | Description | |||| | ELEM\_TYPE | double | Element precision: double or float | | WITH\_GPU | auto | Device for tuned parameters (PVC, A100, ...). Fallback: all CSVs |

### Benchmark Driver

```
./acc_bench.x [nrepeat [batchsize [M [N [K [nc [na [nb]]]]]]]]
```

The kernel shape can also be given as MxNxK:

```
./acc_bench.x 5 30000 13x5x7
```

The first argument can be a file with one set of parameters per line. The batchsize argument accepts K/M/G suffixes for memory-budget mode.

| Argument | Default | Description | |||| | nrepeat | 66 (3+CHECK) | Number of timed repetitions | | batchsize | 30000 | Matrix products per batch (stack size) | | M | 23 | Rows of A and C | | N | M | Columns of B and C | | K | M | Columns of A / rows of B | | nc | batchsize/16 | Number of unique C-matrices | | na | 10nc / Number of unique A-matrices | | nb | 10nc | Number of unique B-matrices |

**Environment Variables (Benchmark)** | Variable | Default | Description | |||| | CHECK | -1 | Accuracy: negative=auto-threshold, 0=off, positive=custom | | CHECK\_H2D | - | Minimum H2D bandwidth (GB/s); fail if below | | CHECK\_DEV | - | Minimum device GFLOPS/s; fail if below | | CHECK\_HST | - | Minimum host GFLOPS/s; fail if below | | DEVICE | 0 | Device index (multi-rank: device = rank % ndevices) | | NREPEAT\_H2D | 1 | H2D copy repetitions for bandwidth measurement | | NREPEAT\_SMM | 1 | SMM kernel launches per timed iteration (for profiling) | | BATCHSIZE\_SMM | - | Override batchsize (and optionally nrepeat: bs,nrep) |

## LIBSMM Kernel Parameters

OpenCL kernels are generated at runtime for any (M, N, K) shape within MAX\_KERNEL\_DIM. Tuned parameters can be embedded at build time or loaded at runtime for better performance.

### Environment Variables (LIBSMM) Transpose kernel:

| Variable                      | Default | Description                                   |
|-------------------------------|---------|-----------------------------------------------|
| OPENCL_LIBSMM_TRANS_BUILDOPTS | -       | Extra OpenCL build options                    |
| OPENCL_LIBSMM_TRANS_INPLACE   | 0       | Non-zero: in-place transpose (no LDS)         |
| OPENCL_LIBSMM_TRANS_BM        | auto    | Block size in M-direction ( $0 < BM \leq M$ ) |

Multiply kernel:

| Variable                    | Default | Description                                   |
|-----------------------------|---------|-----------------------------------------------|
| OPENCL_LIBSMM_SMM_BUILDOPTS | -       | Extra OpenCL build options                    |
| OPENCL_LIBSMM_SMM_PARAMS    | auto    | Tuned parameters: 0=disable, or CSV path      |
| OPENCL_LIBSMM_SMM_BS        | auto    | Intra-kernel mini-batchsize                   |
| OPENCL_LIBSMM_SMM_BM        | auto    | Block size in M-direction ( $0 < BM \leq M$ ) |
| OPENCL_LIBSMM_SMM_BN        | auto    | Block size in N-direction ( $0 < BN \leq N$ ) |
| OPENCL_LIBSMM_SMM_AP        | auto    | Access pattern for parameter/stack array      |
| OPENCL_LIBSMM_SMM_AA        | auto    | Access pattern for A-matrix array             |
| OPENCL_LIBSMM_SMM_AB        | auto    | Access pattern for B-matrix array             |
| OPENCL_LIBSMM_SMM_AC        | auto    | Access pattern for C-matrix array             |

The full list of tunable parameters is available via `tune_multiply.py --help`. Some parameters produce distinct code paths (e.g., SMM\_BS=1 vs SMM\_BS=2), so the parameter space is non-smooth.

## Tuned Parameters

Pre-tuned parameter sets in `params/`:

| File                    | Device                          |
|-------------------------|---------------------------------|
| tune_multiply_PVC.csv   | Intel Data Center GPU Max (PVC) |
| tune_multiply_BMG.csv   | Intel Battlemage (BMG)          |
| tune_multiply_A100.csv  | NVIDIA A100                     |
| tune_multiply_H100.csv  | NVIDIA H100                     |
| tune_multiply_GH200.csv | NVIDIA GH200                    |
| tune_multiply_V100.csv  | NVIDIA V100                     |
| tune_multiply_P100.csv  | NVIDIA P100                     |
| tune_multiply_Mi250.csv | AMD MI250                       |

Parameters are matched by device ID at runtime with best-match fallback. Single and double precision coexist in the same CSV.

```
## Use embedded parameters (default)
```

```
./acc_bench.x 5 30000 13 5 7
```

```
## Disable tuned parameters
```

```
OPENCL_LIBSMM_SMM_PARAMS=0 ./acc_bench.x 5 30000 13 5 7
```

```
## Load from a specific CSV file
```

```
OPENCL_LIBSMM_SMM_PARAMS=params/tune_multiply_PVC.csv ./acc_bench.x 5 30000 13 5 7
```

Rebuild with a specific device's parameters embedded:

```
make realclean
```

```
make WITH_GPU=PVC
```

## Auto Tuning

Uses OpenTuner to explore the parameter space per (M, N, K) kernel. The tuner communicates with the benchmark driver through environment variables.

```
cd samples/smm
pip install -r requirements.txt
```

Build `acc_bench.x` before tuning. Keep GPU clocks and driver state as stable as practical during a run.

**tune\_multiply.py** Use `tune_multiply.py` when you want direct control over one tuning run or over JSON maintenance.

Tune one kernel and write JSON results in the current directory:

```
./tune_multiply.py 13x5x7
```

Limit the search time, choose the JSON directory, and set the benchmark batch size:

```
mkdir -p params/local
./tune_multiply.py 13x5x7 --stop-after=300 -p params/local -s 30000
```

Tune several explicit kernels from a file, one  $M \times N \times K$  per line:

```
printf '%s\n' 13x5x7 23x23x23 32x32x32 > kernels.txt
./tune_multiply.py kernels.txt --stop-after=180 -p params/local
```

Merge JSON files into a CSV file:

```
./tune_multiply.py -m -p params/local -o tune_multiply_local.csv
```

Update stored device names after rebuilding for a different target:

```
./tune_multiply.py -u -p params/local
```

Check existing JSONs without re-tuning:

```
./tune_multiply.py -c -p params/local
```

Useful options:

|        |         |  |               |                                 |  |         |                                     |  |         |                                              |  |          |                                              |  |    |                                  |  |         |                 |  |             |                          |  |              |                       |  |    |                                             |  |
|--------|---------|--|---------------|---------------------------------|--|---------|-------------------------------------|--|---------|----------------------------------------------|--|----------|----------------------------------------------|--|----|----------------------------------|--|---------|-----------------|--|-------------|--------------------------|--|--------------|-----------------------|--|----|---------------------------------------------|--|
| Option | Meaning |  | -stop-after N | Stop the search after N seconds |  | -p path | Directory for JSON input and output |  | -s size | Benchmark batch size, also called stack size |  | -a level | Tuning level: 0=all, 1=most, 2=some, 3=least |  | -m | Merge JSON files into a CSV file |  | -o file | CSV output file |  | -u [device] | Update JSON device names |  | -c [epsilon] | Validate JSON entries |  | -d | Delete outperformed duplicates during merge |  |
|--------|---------|--|---------------|---------------------------------|--|---------|-------------------------------------|--|---------|----------------------------------------------|--|----------|----------------------------------------------|--|----|----------------------------------|--|---------|-----------------|--|-------------|--------------------------|--|--------------|-----------------------|--|----|---------------------------------------------|--|

The tuner can run under MPI. It detects local MPI rank variables and uses them to select `LIBXSTREAM_DEVICE`, so ranks on the same node can use different devices. This is most useful when each rank is tuning a different kernel.

```
mpirun \
  ./tune_multiply.py 13x5x7 --stop-after=300 -p params/local : \
  ./tune_multiply.py 23x23x23 --stop-after=300 -p params/local
```

Interrupted tuning writes the best result seen so far. If you tune a different kernel with `tune_multiply.py` directly, remove any old OpenTuner database first:

```
rm -rf opentuner.db
```

The wrapper below does this cleanup automatically.



**tune\_multiply.sh** Use `tune_multiply.sh` when you want to expand triplet specifications, retune a directory, or split a larger tuning job into parts.

Tune a compact Cartesian-product specification:

```
./tune_multiply.sh -t 300 4 10 15, 6 7 8, 23
```

Triplets are comma-separated groups. The command above expands to all  $M \times N \times K$  combinations with  $M$  in 4 10 15,  $N$  in 6 7 8, and  $K=23$ .

Tune an explicit list instead:

```
./tune_multiply.sh -t 300 1x1x1 2x2x2 13x5x7 23x23x23
```

Split the same work into four parts and run the second part:

```
./tune_multiply.sh -t 300 -j 4 -i 2 4 10 15, 6 7 8, 23
```

Retune every JSON file found in a directory:

```
./tune_multiply.sh -u -p params/local -t 300
```

Limit the generated work before splitting it:

```
./tune_multiply.sh -t 180 -r 4 32 -m 64 -n 100 \  
 4 8 16 32, 4 8 16 32, 4 8 16 32
```

Useful options:

| Option | Meaning | ||| | -t seconds | Time limit per kernel | | -p path | Directory for JSON files | | -s size | Benchmark batch size, also called stack size | | -a level | Tuning level: 0=all, 1=most, 2=some, 3=least | | -u | Retune JSON files found under -p | | -d | Ask the merge step to delete outperformed JSONs | | -c | Continue with the next kernel after an error | | -b | Tune triplets in reverse order | | -j parts | Total number of tuning parts | | -i index | Part to run, using 1-based numbering | | -r low high | Keep kernels with low  $3 < MNK \leq \text{high}3$  | | -m extent | Keep kernels with  $M$ ,  $N$ , and  $K$  no larger than this | | -n count | Keep only the first count kernels | | -f file | Read  $M \times N \times K$  list from a file (one per line) | | -k id | Use a predefined triplet set |

**Tuning from a File** A text file with one  $M \times N \times K$  per line can drive the tuning session. Lines starting with `#` are comments, and inline comments after `#` are stripped. Whitespace within an entry is ignored:

```
./tune_multiply.sh -t 300 -f retune_shapes.txt -p params/local
```

Under MPI, the file entries are partitioned across ranks as usual:

```
mpirun -np 8 ./tune_multiply.sh -t 300 -f retune_shapes.txt -p params/local
```

**Bulk Tuning with MPI** Under MPI, the wrapper defaults `-j` to the MPI world size and `-i` to `rank + 1`. It also forwards a normalized local rank to `tune_multiply.py`, so the Python tuner can select a different device per local rank.

For most MPI launchers, this is enough:

```
mpirun -np 8 ./tune_multiply.sh -t 300 -p params/local \  
 4 10 15, 6 7 8, 23
```

You can still spell out parts explicitly when the launcher or scheduler needs it:

```
mpirun \  
  ./tune_multiply.sh -t 300 -j 4 -i 1 -p params/local \  
    4 10 15, 6 7 8, 23 : \  
  ./tune_multiply.sh -t 300 -j 4 -i 2 -p params/local \  
    4 10 15, 6 7 8, 23 : \  
  ./tune_multiply.sh -t 300 -j 4 -i 3 -p params/local \  
    4 10 15, 6 7 8, 23 : \  
  ./tune_multiply.sh -t 300 -j 4 -i 4 -p params/local \  
    4 10 15, 6 7 8, 23 \  
>out.log 2>&1
```

To retune an existing JSON directory across eight MPI ranks:

```
mpirun -np 8 ./tune_multiply.sh -u -p params/local -t 300
```

**Managing Tuned Parameters** JSON files are the working format. CSV files are the deployable format. A typical retune flow is:

```
make realclean
make WITH_GPU=P100
mkdir -p params/p100
./tune_multiply.sh -u -p params/p100 -t 300
./tune_multiply.py -m -p params/p100 -o params/tune_multiply_P100.csv
```

Keep GPU driver state persistent during tuning (e.g., `nvidia-smi -pm ENABLED` on headless NVIDIA systems).

## Stencil

Finite-difference stencil computation for seismic wave propagation (RTM, FWI) on GPUs. Three kernel paths are available:

- **FP32** (default): dedicated scalar kernel with SLM blocking
- **BF16-DPAS** (`STENCIL_BF16=1`): Dekker-split BF16 via hardware matrix units (DPAS/XMX)
- **INT8-DPAS** (`STENCIL_INT8=1`): Ozaki-1 INT8 with carried-forward exponent

The DPAS paths achieve single-precision accuracy from low-precision tensor operations.

## Method

The 3D Laplacian is decomposed into three 1D operators applied along each axis (dimension splitting). Each 1D operator D is a BLK x BLK banded Toeplitz matrix (bandwidth  $2R+1$ ) representing the FD stencil.

Parameters (compile-time):

```
BLK          = 32      block side length (32^3 output cube)
RADIUS       = 4       half-order (8th-order FD, 9-point stencil)
```

**FP32 path (default)** The dedicated FP32 kernel uses SLM tiling with a sliding window along the slow axis. Work-group size 32x8, cooperative SLM fill, no DPAS dependency. Supports XYZ and ZYX memory layouts, PML absorbing boundaries, and compact/dispersion-fitted operator methods.

**BF16 path (`STENCIL_BF16=1`)** The wavefield block P and operator D are Dekker-split into BF16 digits. The matrix product  $D \times P$  is computed as a sum of BF16 x BF16 DPAS calls that accumulate into FP32.

```
NDIGITS_A = 2      BF16 digits for operator (11-bit range)
NDIGITS_X = 3      BF16 digits for wavefield
```

Products per dimension:  $NDIGITS\_A \times NDIGITS\_X = 6$  DPAS calls. Isotropic total:  $3 \times 6 = 18$  DPAS calls per output block.

**INT8 path (`STENCIL_INT8=1`)** Ozaki-1 slicing: the FD operator fits in a single 7-bit signed digit (`NSLICES_A=1`). The wavefield is sliced into 1-3 digits at runtime depending on local dynamic range.

```
NSLICES_A = 1      INT8 digits for operator
NSLICES_X = 1-3    INT8 digits for wavefield (runtime adaptive)
K_PAD_I8  = 64     K dimension (k=32 DPAS alignment)
```

Products per dimension:  $1 \times nslices\_eff = 1-3$  DPAS calls. Isotropic total:  $3 \times (1-3) = 3-9$  DPAS calls per output block.

A per-block carried-forward exponent (`exp_buf`) tracks the dynamic range across time steps. Double-buffered output-based tracking avoids race conditions between neighboring blocks.

## 2D Block I/O

The kernel exploits Intel 2D block load instructions for the A-side (operator) and sub-group block reads for the B-side (wavefield in SLM):

BF16 path:

```
A: intel_sub_group_2d_block_read_16b_8r16x1c
B: intel_sub_group_2d_block_read_transform_16b_16r16x1c (VNNI)
```

INT8 path:

```
A: intel_sub_group_2d_block_read_8b_8r32x1c (k=32)
B: intel_sub_group_block_read8 from row-major SLM
```

The operator D is stored dense (banded zeros included). The structural zeros cost no memory bandwidth (D fits in L1 after first access) and no effective compute (kernel is memory-bound on wavefield reads).

The INT8 kernel gathers the wavefield directly into SLM during the fill loop (no separate preprocess kernel).

## Build

```
make DBG=1 PEDANTIC=2
```

Requires LIBXS (sibling directory ../../libxs) and an OpenCL 2.0 runtime with `cl_intel_subgroup_2d_block_io` and DPAS support.

Build-time defines go into ECFLAGS, e.g. `make ECFLAGS="-DSTENCIL_PML_W=27"` (DFLAGS and CFLAGS are assembled by the build and a command line assignment would replace them):

```
STENCIL_PML_W      absorbing layer width in cells (default: 20). A
                   caller supplying its own eta sets its own damping
                   width here; a larger value is safe, a smaller one
                   leaves damped cells unattenuated.
```

```
make OCL=0
```

Compiles the FP32 kernel (kernels/stencil\_fp32.cl) with the C compiler and runs it on the CPU instead of the GPU. Needs neither OpenCL nor the LIBXSTREAM library. The kernel is specialized at build time, so the work-group shape and optionally the grid extent are pinned by CPUDEF:

```
make OCL=0 CPUDEF="-DWG_X=128 -DWG_Y=8"
```

WG\_X defaults to the full fast axis (capacity 1024) and WG\_Y to 8. Bind the threads, which is what `make OCL=0 test` does:

```
OMP_PROC_BIND=spread OMP_PLACES=cores ./stencil.x -n 512
```

The kernel is compiled once per combination the run may ask for, and `stencil_configure` selects between them, which is how these environment variables keep working without a JIT:

```
STENCIL_BF16S      wavefield storage limbs (1 = single BF16, 2 = two)
STENCIL_FP16S      IEEE FP16 wavefield storage (0/1)
STENCIL_HALO       halo padding size per axis
STENCIL_METHOD     operator method (0-3)
STENCIL_LAYOUT     accepted when it names the layout built in
STENCIL_PML        accepted when it matches the build
```

Four storage formats times four operator radii times the two boundary treatments (clamp the gather, or read it out of the halo, which the grid decides as on the device) are 32 instances, about 260 KB of code. Drop the compact radii for a build of eight:

```
make OCL=0 CPUDEF="-DSTENCIL_CPU_COMPACT=0"
```

Asking for a combination the build does not carry is an error naming the flag to rebuild without, never a silent substitution. `STENCIL_PADDED` forces the boundary treatment and then compiles only that half.

The benchmark driver reaches the host build through `--exe`, which makes the CPU kernel sweepable like any device variant:

```
make OCL=0 BLDDIR=/tmp/b/obj OUTDIR=/tmp/b
./stencil.py --exe /tmp/b/stencil.x --kernel bf16s --sizes 256,512
```

BF16S costs about 8% against FP32 and FP16S about 3x, because the FP16 conversion does not vectorize in the C89 mode `PEDANTIC=2` selects. `BF16S=2` (`--kernel bf16s-split`) holds two limbs and therefore moves as many bytes as FP32: it buys precision, not bandwidth.

## Usage

```
./stencil.x [options]
```

```
-n <N>          grid dimension (NxNxN, default 256)
-nx/ny/nz <N>   individual grid dimensions
-t <steps>       time steps (default 100)
-d <dims>        operator terms: 3=isotropic, 9=TTI (default 3)
-h <spacing>     grid spacing in meters (default 10.0)
-v <model>       velocity: const | grad | layered | <file.bin>
-vmin <vel>      minimum velocity m/s (default 1500)
-vmax <vel>      maximum velocity m/s (default 4500)
-f <freq>        source frequency Hz (default 25)
-w <steps>       warmup steps excluded from timing (default 5)

-seg-salt        preset: SEG/EAGE Salt (676x676x210, h=20m)
-overthrust      preset: SEG/EAGE Overthrust (801x801x187, h=25m)
```

Performance is reported as GPoints/s (grid points updated per second).

## Velocity Models

The driver supports loading external velocity models as flat binary files (float32, x-fastest order). Several standard models are publicly available for benchmarking.

```
Grid:    676 x 676 x 210    (n1=210 depth, n2=n3=676 lateral)
Spacing: 20 m
Range:    1500 - 4482 m/s
Size:     ~385 MB (float32)
Notes:    Salt body with sharp velocity contrasts.
          Standard RTM benchmark (Aminzadeh et al., 1997).
```

Original (20 m): <https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Salt/fvp>

Resampled (40 m, 115x338x338, for cheaper runs): [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Salt/BENCH\\_27PT/v.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Salt/BENCH_27PT/v.bin)

Reference wavefield (frequency-domain CBS solution): [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Salt/BENCH\\_27PT/cbs\\_salt\\_real.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Salt/BENCH_27PT/cbs_salt_real.bin) [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Salt/BENCH\\_27PT/cbs\\_salt\\_imag.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Salt/BENCH_27PT/cbs_salt_imag.bin)

Grid: 801 x 801 x 187 (n1=187 depth, n2=n3=801 lateral)  
Spacing: 25 m  
Range: 2179 - 6000 m/s  
Size: ~450 MB (float32)  
Notes: Layered geology with thrust faults. Smoother than Salt.  
Reference: Aminzadeh, Brac, and Kunz (1997).

Original (25 m): <https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/fvp>

Resampled (50 m, 94x401x401): <https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/fvp.r50>

2D section (slice i3=455, 187x801): <https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/overthrust2D>

Reference wavefield (CBS, 50 m grid): [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/BENCH\\_27PT/v.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/BENCH_27PT/v.bin) [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/BENCH\\_27PT/cbs\\_overthrust\\_real.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/BENCH_27PT/cbs_overthrust_real.bin) [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/BENCH\\_27PT/cbs\\_overthrust\\_imag.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/BENCH_27PT/cbs_overthrust_imag.bin)

Grid: 2301 x 751 (n1=751 depth, n2=2301 lateral)  
Spacing: 4 m  
Notes: Classic 2D migration benchmark (Versteeg, 1994).

Velocity (vp.bin): <https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Marmousi/vp.bin>

Density (rho.bin): <https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Marmousi/rho.bin>

Grid: 1911 x 12601  
Spacing: 6.25 m  
Notes: Challenging sub-salt imaging (Billette and Brandsberg-Dahl, 2004).  
Available from SEG open data.

SEG wiki page (registration may be required): [https://wiki.seg.org/wiki/2004\\_BP\\_Velocity\\_Estimation\\_Benchmark\\_Model](https://wiki.seg.org/wiki/2004_BP_Velocity_Estimation_Benchmark_Model)

Grid: 641 x 209 (n1=209 depth, n2=641 lateral)  
Spacing: 25 m  
Type: Acoustic VTI (vertical transverse isotropy)  
Fields: Vp, epsilon, delta, eta, Vnmo, density  
Notes: Representative of North Sea environments.  
Exercises anisotropic operator (pure terms with modified coefficients; cross-terms vanish for VTI).

Download (Geoazur WIND): [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/vp\\_true.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/vp_true.bin) [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/epsilon\\_true.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/epsilon_true.bin) [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/delta\\_true.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/delta_true.bin) [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/eta\\_true.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/eta_true.bin) [https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/rho\\_true.bin](https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Valhall2D/rho_true.bin)

Grid: ~1911 x 12480  
Spacing: 6.25 m  
Type: Acoustic TTI (tilted transverse isotropy)  
Fields: Vp, epsilon, delta, theta (tilt angle)  
Notes: Canonical TTI benchmark (Shafiq et al., 2007).  
Exercises full cross-derivative kernel with non-zero  
tilt angles. This is the standard reference for TTI  
stencil performance.

SEG wiki page (registration required): [https://wiki.seg.org/wiki/2007\\_BP\\_Anisotropic\\_Velocity\\_Benchmark\\_Model](https://wiki.seg.org/wiki/2007_BP_Anisotropic_Velocity_Benchmark_Model)

**Synthetic 3D TTI (from Overthrust)** No public 3D TTI model exists. The standard practice in GPU stencil papers is to overlay synthetic Thomsen parameters on the 3D Overthrust velocity model:

Vp: from Overthrust (801x801x187, 25 m)  
epsilon: constant 0.1 - 0.2 (or linear gradient)  
delta: constant 0.05 - 0.1  
theta: smooth tilt (e.g., 20-45 degrees, linearly varying with depth)  
phi: azimuthal angle (0 or smooth variation)

This gives a realistic velocity structure with controlled anisotropy parameters for benchmarking the TTI kernel (-d 9 mode).

**Data hosting (Geoazur WIND project)** All models above (except BP 2004) are hosted by the WIND project at Universite Cote d'Azur / Geoazur. Index page:

<https://www.geoazur.fr/WIND/bin/view/Main/Data/>

**Usage with this benchmark** Download the velocity binary and run:

```
wget https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Salt/fvp
./stencil.x -seg-salt -v fvp -t 1000
```

```
wget https://www.geoazur.fr/WIND/pub/nfs/FWI-DATA/GEOMODELS/Overthrust/fvp
./stencil.x -overthrust -v fvp -t 1000
```

The binary format is flat float32 with n1 (depth) as the fastest dimension. Velocities are in m/s; the driver squares them internally. Byte order is native (little-endian on x86).

## Performance Metric

The standard metric for stencil codes is GPoints/s:

$$\text{GPoints/s} = (\text{nx} \times \text{ny} \times \text{nz} \times \text{nsteps}) / (\text{time} \times 1\text{e9})$$

This measures throughput independent of the internal algorithm (direct FD, GEMM-based, FFT, etc.) and is directly comparable across:

- Devito (Imperial College) --- symbolic PDE -> optimized C/OpenCL
- minimod (TotalEnergies) --- open-source GPU stencil mini-app
- Published results on PVC, A770, H100, MI300X

For reference, a memory-bandwidth-bound 8th-order stencil on PVC (HBM bandwidth ~3.2 TB/s) achieves approximately 200-400 GPoints/s depending on grid size and occupancy.

## TTI (Anisotropic) Mode

With -d 9, the kernel computes the full tilted transverse isotropy operator with cross-derivative terms. Each cross-term uses SLM to buffer the intermediate result between two GEMM phases:

```
T  = D_j x P          (first GEMM)
T  = c_ij . T          (point-wise anisotropy scaling)
Y += D_i x T          (second GEMM)
```

SLM budget per cross-term:  $2 \times K\_PAD \times XMX\_N \times \text{sizeof}(\text{ushort})$ , rounded up to the DPAS K block size. Total for 3 cross-terms with barrier: fits within 128 KB (Xe2/BMG).

## Operator Paths

The standard direct kernel applies one wide-reach operator (radius-4, 9-point) per dimension. The compact variants are independent runtime paths: they compile the same Ozaki-split BF16 x BF16 DPAS primitive with a smaller gather radius and rely on time evolution to build longer effective reach.

Methods (selected via STENCIL\_METHOD environment variable):

0 = direct K=1, r=4 standard high-order (default) 1 = compact-r1 K=4, r=1 compact tridiagonal runtime path 2 = compact-r2 K=2, r=2 compact pentadiagonal runtime path 3 = compact-fit K=4, r=1 placeholder for dispersion-fitted coefficients

The “compact-fit” mode uses the free degrees of freedom in the K-factor coefficient product to match or exceed the dispersion quality of the standard 8th-order stencil at target frequencies, while retaining the memory savings of the cascade. Coefficients are precomputed once at initialization for the grid’s frequency content.

The first compact path is implemented for the isotropic apply kernel in methods 1-3. TTI cross-terms still use the direct two-phase DPAS path.

Environment variables controlling kernel specialization:

```
STENCIL_BF16      BF16-DPAS kernel (1=native BF16, 2=FP32-split via BF16)
STENCIL_BF16S     BF16 wavefield storage limbs (0=off, 1=single 2-byte
limb halves traffic, 2=two limbs FP32-equivalent)
STENCIL_FP16S     IEEE FP16 wavefield storage (0/1); single 2-byte value
per point, halves traffic (10-bit mantissa, FP32 kernel)
STENCIL_INT8      INT8-DPAS kernel (1=native INT8, 2=FP32-split via INT8)
STENCIL_METHOD    operator method (0-3, default 0)
STENCIL_STRIPS_PER_WG
    adjacent N-strips handled by one work-group (default: 2)
STENCIL_SG        subgroup size override (default: device preferred)
STENCIL_GRF256    force 256-GRF mode (0/1, default: auto)
STENCIL_TRIM      drop least-significant digit products (BF16 path)
STENCIL_LU        loop unroll strategy (-1=none, 0=inner, 1=outer)
STENCIL_LAYOUT    memory layout (0=XYZ, 1=blocked, 2=ZYX)
STENCIL_HALO      halo/padding size per axis
STENCIL_PML       enable PML absorbing boundary (0/1)
STENCIL_FP32_WG_X, STENCIL_FP32_WG_Y
    FP32 direct-kernel work-group shape (default: 32x8)
    STENCIL_FP32_SBLOCK
        FP32 direct-kernel slow-axis SLM microblock (1 or 2,
        default: 2 except ZYX+PML)
STENCIL_FP32_BLOCK_IO
    enable FP32 Intel 2D block reads for padded 32x8 cases (0/1, default: 0)
```

Specialized kernels are compiled on first use and cached in a thread-safe registry keyed by the method, compact-step parameters, strip grouping, subgroup size, packed specialization flags, FP32 work-group shape, grid shape, and term count. Subsequent launches with the same parameters are zero-cost.

Future extension: per-block adaptive method selection (different K in regions with different velocity/frequency content). This requires per-block dispatch logic and is not yet implemented.

## Integration

The stencil API (stencil\_opencl.h) accepts device pointers for all buffer arguments (wavefield, output, velocity). No host-device transfers happen inside the dispatch path — data stays on-device across the full time-stepping loop.

**Initialization and USM control** LIBXSTREAM provides libxstream\_init\_config for explicit control over the memory model before initialization:

```
libxstream_init_config_t cfg;
libxstream_init_config_default(&cfg); /* all fields = -1 (default) */
cfg.usm = 1; /* force Intel USM extensions */
cfg.device = 0; /* select device index */
libxstream_init_config(&cfg);
```

USM levels (cfg.usm):

```
-1 env/default (LIBXSTREAM_USM, or SVM coarse-grain fallback)
0 disable USM, force clCreateBuffer path
1 Intel USM extensions (clDeviceMemAllocINTEL)
2 OpenCL 2.0 SVM coarse-grain only
3 OpenCL 2.0 SVM with device-reported capabilities
```

When USM is active (level 1), device pointers from any source (SYCL, Level Zero, raw clDeviceMemAllocINTEL) can be passed to the stencil API without conversion.

**SYCL interoperability** SYCL USM device allocations (sycl::malloc\_device) can be passed directly to stencil\_apply\_laplacian on Intel GPUs. The underlying mechanism is clSetKernelArgMemPointerINTEL, which accepts the same raw pointers that SYCL produces via the Level Zero unified runtime.

Requirements:

- Initialize LIBXSTREAM with cfg.usm = 1, or set the environment variable LIBXSTREAM\_USM=1.
- The SYCL queue and the libxstream OpenCL context must target the same physical device (shared Level Zero context).
- Synchronization: use sycl::queue::ext\_oneapi\_submit\_barrier() before calling stencil\_apply\_laplacian, and libxstream\_stream\_sync after, to order work correctly.
- No format conversion needed: float32 USM buffers are used as-is (velocity is expected as  $v^2$ ).

**External OpenCL consumers** Any OpenCL application can call the API by passing cl\_mem-backed device pointers obtained via libxstream\_mem\_allocate, or USM pointers from clDeviceMemAllocINTEL directly. The dispatch layer auto-detects USM availability and selects the appropriate kernel argument-setting path.

**Scalar proxy (non-DPAS devices)** The kernels include a scalar fallback that runs on any OpenCL 1.2+ device without DPAS/XMX hardware. This enables functional testing and correctness verification on iGPUs, CPUs, and other devices. Performance is not representative — the proxy exists purely for development and validation.

## File Layout

|                   |                                                     |
|-------------------|-----------------------------------------------------|
| samples/stencil/  |                                                     |
| Makefile          | build rules                                         |
| README.md         | this file                                           |
| stencil.c         | host driver (benchmark, model loading)              |
| stencil.py        | Python benchmark harness (CSV/plot output)          |
| stencil_opencl.c  | OpenCL context, kernel dispatch, exp_buf management |
| stencil_opencl.h  | public API, compile-time parameters                 |
| stencil_kernels.h | generated at build time from .cl sources            |
| kernels/          |                                                     |
| stencil_common.cl | parameters, gather macros, layout indexing          |
| stencil_fp32.cl   | FP32 path: stencil_apply_direct (default)           |
| stencil_bf16.cl   | BF16 path: stencil_apply, stencil_apply_tti         |
| stencil_int8.cl   | INT8 path: stencil_apply_int8                       |



```
libxstream/opencv/  
  libxstream_common.h  IEEE utilities, BF16 conversion, EXP2I, unroll macros  
  libxstream_ozaki.h   Dekker split, BF16 DPAS macros
```